Student Name: Gowtham Kumar Kamuni
Student ID: A20549435
LAB EXPERIMENT 2

OUTPUTS:

Jetson Orin Nano – Image Classification Output Summary

This document presents the results from image classification tests performed using the GoogLeNet model on the Jetson Orin Nano. Each test image was processed through the imagenet.py script provided in NVIDIA's jetson-inference library.

For each case, the original input image is shown alongside the corresponding output with a predicted label and confidence score overlayed. The intent of this output file is to visually demonstrate the performance, accuracy, and behavior of the classification model across a variety of real-world test cases captured during the lab. All outputs were generated directly on the Jetson device as part of the Lab Experiment 2 demonstration.

1. Preliminary Work:

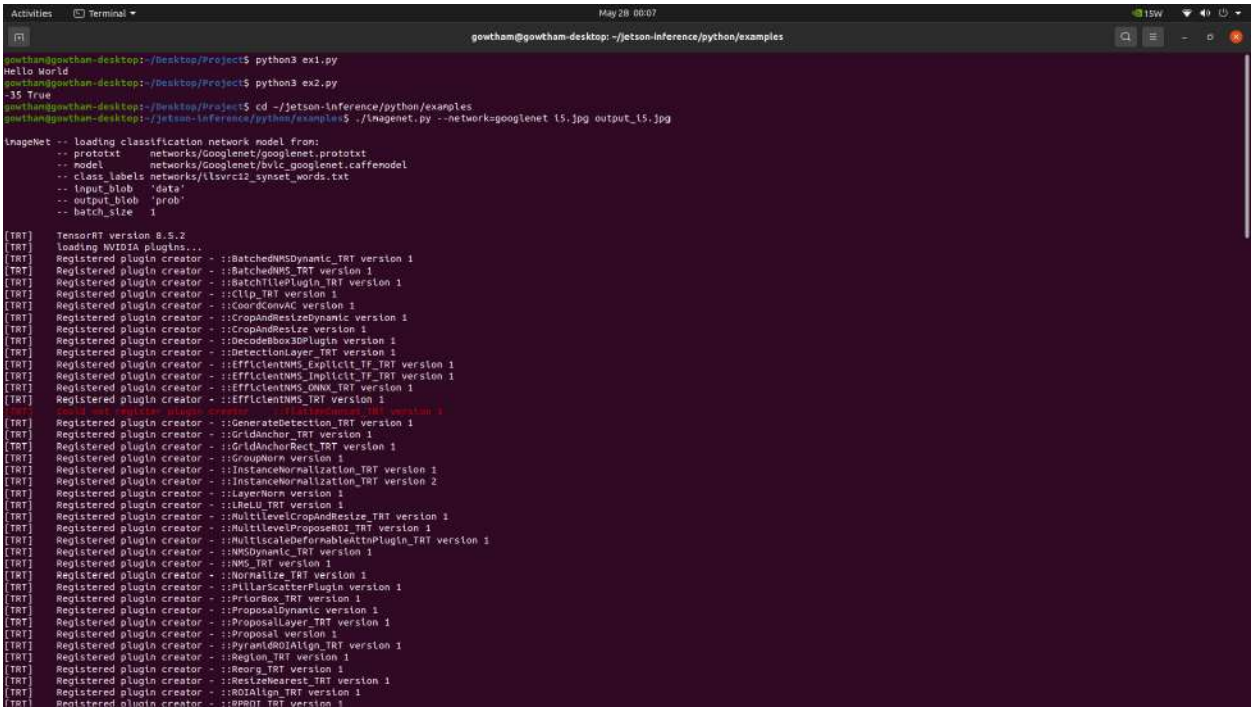

Fig: Preliminary Work – Python Execution on Jetson Orin Nano

The screenshot above shows successful execution of two basic Python programs on Jetson.

- ex1.py prints **“Hello World”**, confirming Python is running correctly
- ex2.py implements a simple **BankAccount class**, showcasing object-oriented programming with a withdrawal operation and overdraw check.
- The terminal also shows navigation to the jetson-inference directory, ready for image classification.
  This validated the Python environment and ensured Jetson was prepared for the main classification tasks.

Test Case 1:

Fig: Timber Wolf

Input Image: A high-resolution image of a timber wolf walking in the snow.

Output Description: The image was successfully recognized with **91.27% confidence** as a "timber wolf". This confirms that the model performs very well on clear, centered images that align with its training dataset.

Test Case 2:

Fig: Backpack

Input: A black “Mountain Adventure” backpack placed on a couch.
Output: Classified as **“blue jean, denim”** with **62.57% confidence**.

Comment: This was a misclassification. The model seemed to rely on color/ texture, likely confusing the lighting on the backpack with denim. Shows limitations in object recognition beyond color similarity.

Test Case 3:

Fig: Laptop

Input: Closed MacBook Air with custom stickers on a brown sofa.
Output: Classified as “laptop, laptop computer” with **18.60%** confidence.
Comment: While the model did identify the object as a laptop, the confidence was quite low. This shows that external visual elements like stickers may impact how clearly the object is recognized by the model.

Test Case 4:

Fig: Fruits and Vegetables

Input: A carrot, a spotted banana, and a tomato placed on a couch.
Output: Classified as **"orange"** with **23.63% confidence**.
Comment: Partial misclassification. While none of the items were oranges, the dominant color may have influenced the prediction. Demonstrates the model's reliance on color features more than object structure in some cases.

Test Case 5:

Fig: Skateboard

Input: A skateboard on a carpeted floor next to a wall.
Output: Classified as **“coho salmon, silver salmon”** with **14.53% confidence**.
Comment: Incorrect result. The skateboard was completely misclassified, likely due to shape and texture mismatch in the model’s training set. This highlights GoogLeNet’s limitations with unfamiliar objects.

<u>Test Case 6:</u>


Fig: Friend

Input Image: A portrait photo of a person (Rajesh) sitting at a desk in a casual indoor setup.
Output Description: The model classified the image with **40.56% confidence** as "miniskirt, mini", which is a completely inaccurate result. This highlights a limitation of the model, it misclassifies images that contain people in natural poses, especially when their clothing doesn't match clear class labels in the training set. This test shows the importance of dataset diversity and the limits of generalization in pre-trained models.

The following screenshots display the terminal output generated during image classification using the Jetson Orin Nano and the imagenet.py script. Each test case (from i1.jpg to i5.jpg) was processed using the GoogLeNet model. The outputs show the classification labels, confidence scores, image processing pipeline initialization, and detailed inference timing reports from both the CPU and CUDA. These logs confirm successful execution and demonstrate the internal workings of the model during inference.

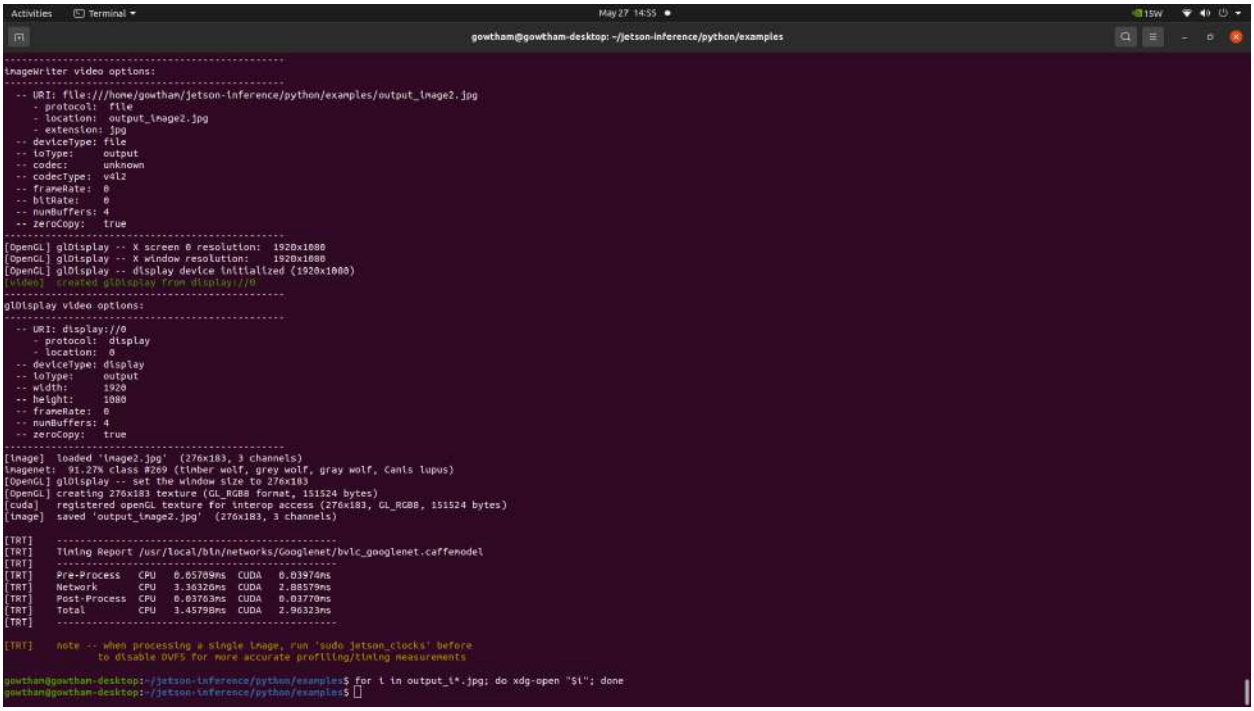

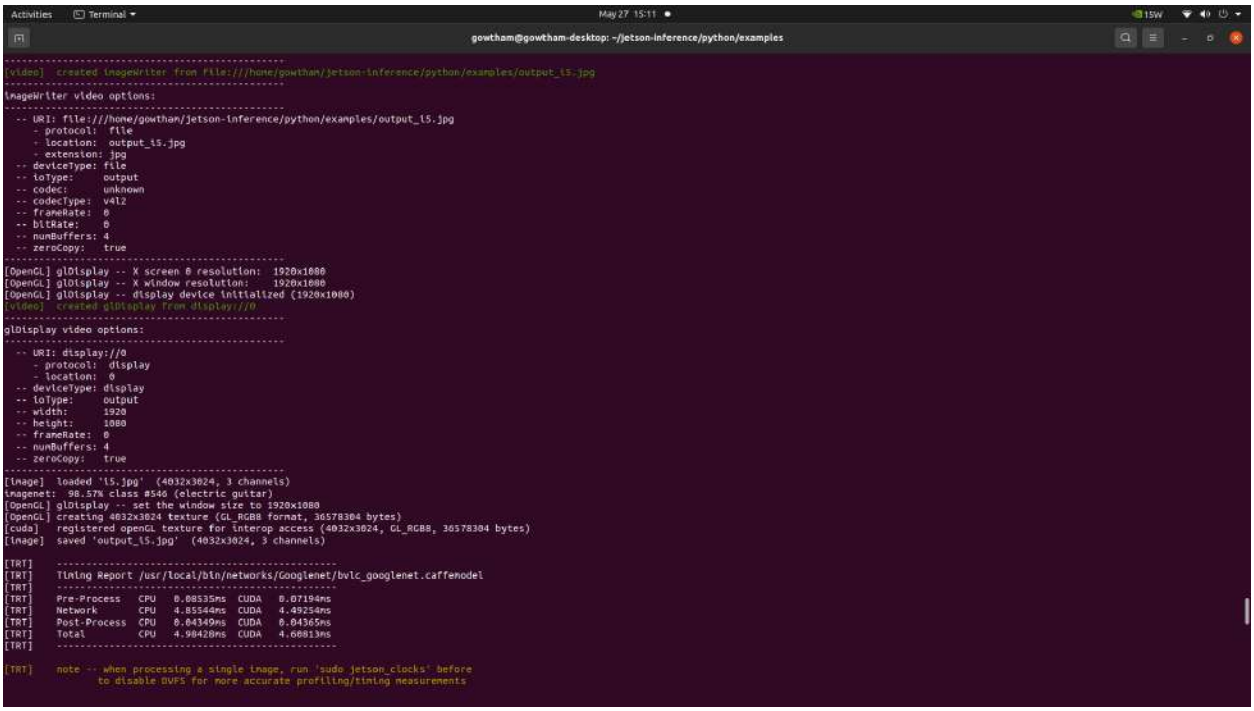

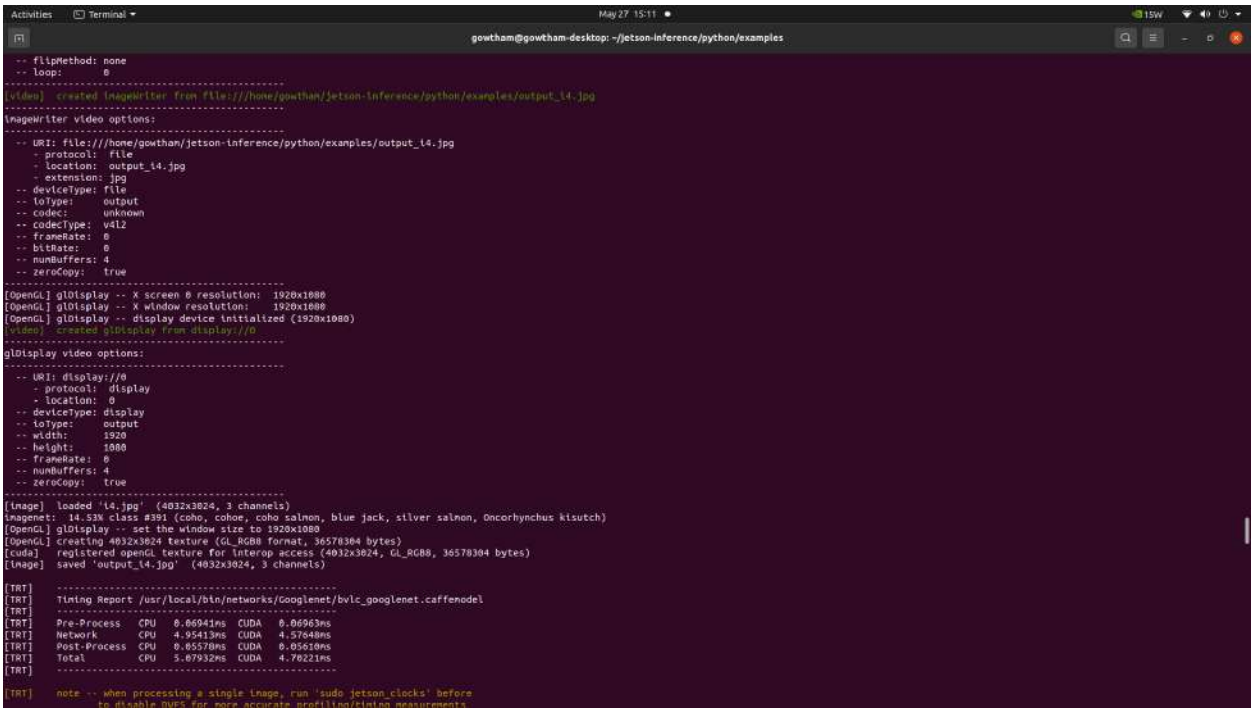
Activities
Terminal
May 27 15:11
gowtham@gowtham-desktop: ~/jetson-inference/python/examples
-- flipMethod: none
-- loop: 0
[video] created imageWriter from file:///home/gowtham/jetson-inference/python/examples/output_l4.jpg
imageWriter video options:
-- URI: file:///home/gowtham/jetson-inference/python/examples/output_l4.jpg
- protocol: file
- location: output_l4.jpg
- extension: jpg
-- deviceType: file
-- ioType: output
-- codec: unknown
-- codecType: v4l2
-- frameRate: 0
-- bitRate: 0
-- numBuffers: 4
-- zeroCopy: true
[OpenGL] glDisplay -- X screen 0 resolution: 1920x1080
[OpenGL] glDisplay -- X window resolution: 1920x1080
[OpenGL] glDisplay -- display device initialized (1920x1080)
[video] created glDisplay from display://0
glDisplay video options:
-- URI: display://0
- protocol: display
- location: 0
-- deviceType: display
-- ioType: output
-- width: 1920
-- height: 1080
-- frameRate: 0
-- numBuffers: 4
-- zeroCopy: true
[image] loaded 'l4.jpg' (4032x3024, 3 channels)
imagenet: 14.53% class #391 (coho, cohoe, coho salmon, blue jack, silver salmon, Oncorhynchus kisutch)
[OpenGL] glDisplay -- set the window size to 1920x1080
[OpenGL] creating 4032x3024 texture (GL_RGB8 format, 36578304 bytes)
[cuda] registered openGL texture for interop access (4032x3024, GL_RGB8, 36578304 bytes)
[image] saved 'output_l4.jpg' (4032x3024, 3 channels)
[TRT]
[TRT] Timing Report /usr/local/bin/networks/Googlenet/bvlc_googlenet.caffemodel
[TRT]
[TRT] Pre-Process CPU 0.06941ms CUDA 0.06963ms
[TRT] Network CPU 4.95413ms CUDA 4.57640ms
[TRT] Post-Process CPU 0.85578ms CUDA 0.05610ms
[TRT] Total CPU 5.87932ms CUDA 4.70221ms
[TRT]
[TRT] note -- when processing a single image, run 'sudo jetson_clocks' before

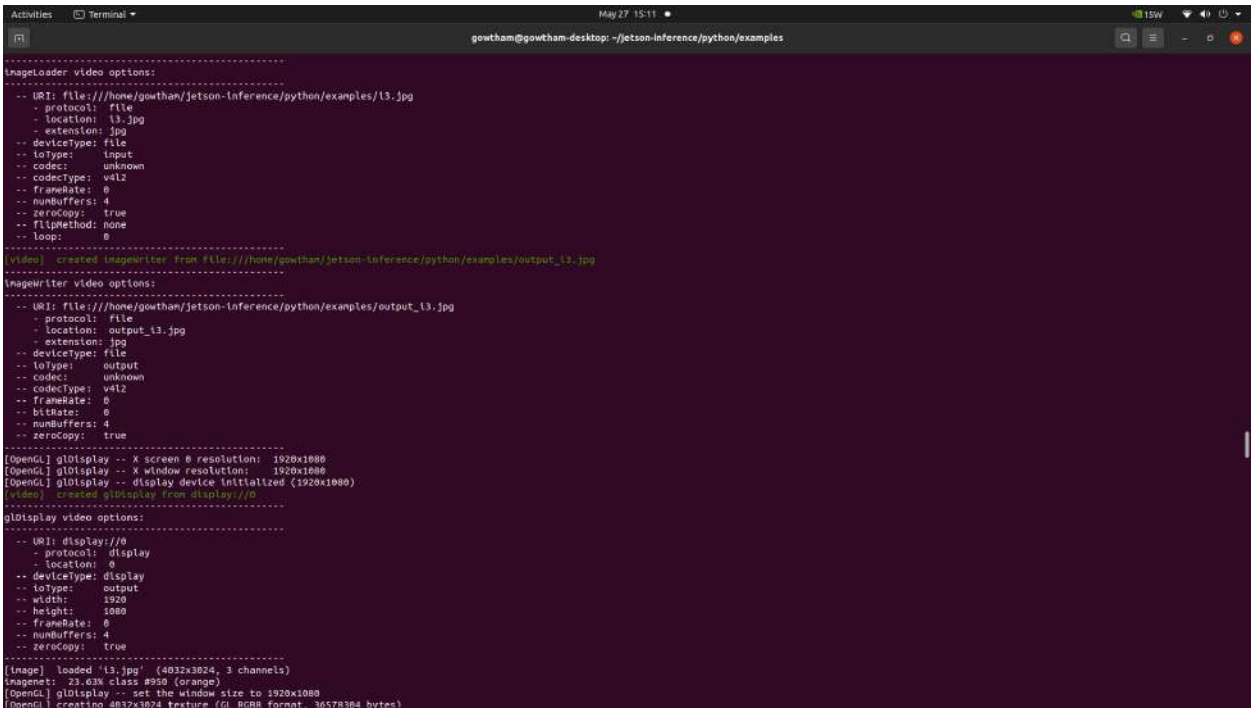
Activities
Terminal
May 27 15:11
gowtham@gowtham-desktop: ~/jetson-inference/python/examples
imageLoader video options:
-- URI: file:///home/gowtham/jetson-inference/python/examples/l3.jpg
- protocol: file
- location: l3.jpg
- extension: jpg
-- deviceType: file
-- ioType: input
-- codec: unknown
-- codecType: v4l2
-- frameRate: 0
-- numBuffers: 4
-- zeroCopy: true
-- flipMethod: none
-- loop: 0
[video] created imageWriter from file:///home/gowtham/jetson-inference/python/examples/output_l3.jpg
imageWriter video options:
-- URI: file:///home/gowtham/jetson-inference/python/examples/output_l3.jpg
- protocol: file
- location: output_l3.jpg
- extension: jpg
-- deviceType: file
-- ioType: output
-- codec: unknown
-- codecType: v4l2
-- frameRate: 0
-- bitRate: 0
-- numBuffers: 4
-- zeroCopy: true
[OpenGL] glDisplay -- X screen 0 resolution: 1920x1080
[OpenGL] glDisplay -- X window resolution: 1920x1080
[OpenGL] glDisplay -- display device initialized (1920x1080)
[video] created glDisplay from display://0
glDisplay video options:
-- URI: display://0
- protocol: display
- location: 0
-- deviceType: display
-- ioType: output
-- width: 1920
-- height: 1080
-- frameRate: 0
-- numBuffers: 4
-- zeroCopy: true
[image] loaded 'l3.jpg' (4032x3024, 3 channels)
imagenet: 23.63% class #950 (orange)
[OpenGL] glDisplay -- set the window size to 1920x1080

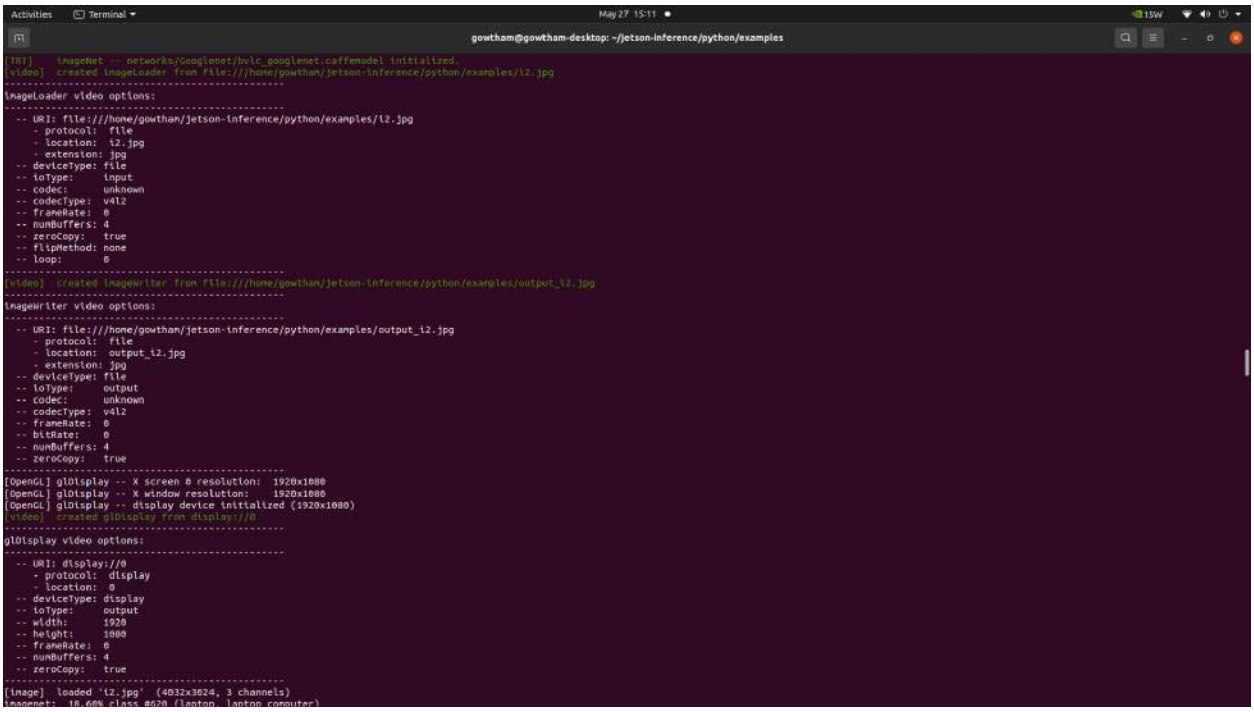
Activities
Terminal
May 27 15:11
15W
gowtham@gowtham-desktop: ~/jetson-inference/python/examples
[TRT]   imageNet -- networks/Googlenet/bvlc_googlenet.caffemodel initialized.
[video] created imageLoader from file:///home/gowtham/jetson-inference/python/examples/i2.jpg
imageLoader video options:
-- URI: file:///home/gowtham/jetson-inference/python/examples/i2.jpg
- protocol: file
- location: i2.jpg
- extension: jpg
-- deviceType: file
-- ioType: input
-- codec: unknown
-- codecType: v4l2
-- frameRate: 0
-- numBuffers: 4
-- zeroCopy: true
-- flipMethod: none
-- loop: 0
[video] created imageWriter from file:///home/gowtham/jetson-inference/python/examples/output_i2.jpg
imageWriter video options:
-- URI: file:///home/gowtham/jetson-inference/python/examples/output_i2.jpg
- protocol: file
- location: output_i2.jpg
- extension: jpg
-- deviceType: file
-- ioType: output
-- codec: unknown
-- codecType: v4l2
-- frameRate: 0
-- bitRate: 0
-- numBuffers: 4
-- zeroCopy: true
[OpenGL] glDisplay -- X screen 0 resolution: 1920x1080
[OpenGL] glDisplay -- X window resolution: 1920x1080
[OpenGL] glDisplay -- display device initialized (1920x1080)
[video] created glDisplay from display://0
glDisplay video options:
-- URI: display://0
- protocol: display
- location: 0
-- deviceType: display
-- ioType: output
-- width: 1920
-- height: 1080
-- frameRate: 0
-- numBuffers: 4
-- zeroCopy: true
[image] loaded 'i2.jpg' (4032x3024, 3 channels)
imagenet: 18.60% class #620 (laptop, laptop computer)

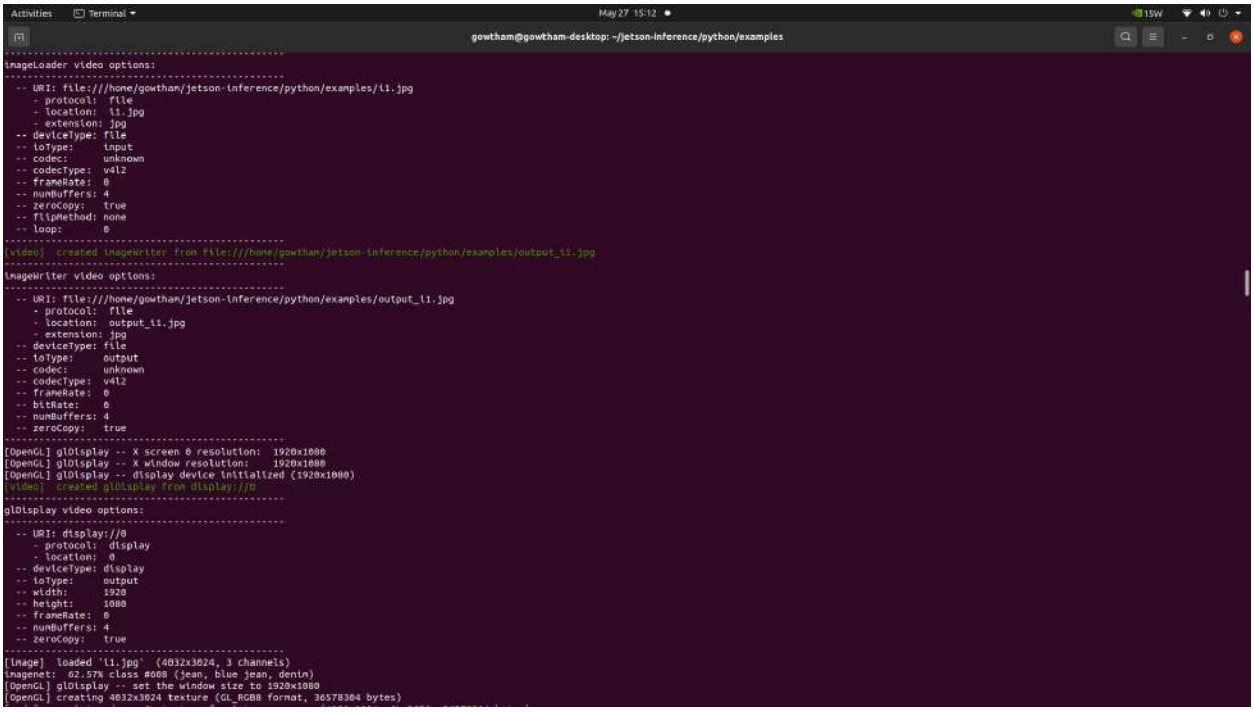
Activities
Terminal
May 27 15:12
15W
gowtham@gowtham-desktop: ~/jetson-inference/python/examples
imageLoader video options:
-- URI: file:///home/gowtham/jetson-inference/python/examples/i1.jpg
- protocol: file
- location: i1.jpg
- extension: jpg
-- deviceType: file
-- ioType: input
-- codec: unknown
-- codecType: v4l2
-- frameRate: 0
-- numBuffers: 4
-- zeroCopy: true
-- flipMethod: none
-- loop: 0
[video] created imageWriter from file:///home/gowtham/jetson-inference/python/examples/output_i1.jpg
imageWriter video options:
-- URI: file:///home/gowtham/jetson-inference/python/examples/output_i1.jpg
- protocol: file
- location: output_i1.jpg
- extension: jpg
-- deviceType: file
-- ioType: output
-- codec: unknown
-- codecType: v4l2
-- frameRate: 0
-- bitRate: 0
-- numBuffers: 4
-- zeroCopy: true
[OpenGL] glDisplay -- X screen 0 resolution: 1920x1080
[OpenGL] glDisplay -- X window resolution: 1920x1080
[OpenGL] glDisplay -- display device initialized (1920x1080)
[video] created glDisplay from display://0
glDisplay video options:
-- URI: display://0
- protocol: display
- location: 0
-- deviceType: display
-- ioType: output
-- width: 1920
-- height: 1080
-- frameRate: 0
-- numBuffers: 4
-- zeroCopy: true
[image] loaded 'i1.jpg' (4032x3024, 3 channels)
imagenet: 62.57% class #608 (jean, blue jean, denim)
[OpenGL] glDisplay -- set the window size to 1920x1080
[OpenGL] creating 4032x3024 texture (GL_RGB8 format, 36578304 bytes)